

PERSONAL · AI · FOREVER

AI of One's Own

Their intelligence.

ChatGPT, Claude + *your memory.*

BY NOBORU IKUTA

An AI that remembers you. *Forever.*

Everywhere you go. On every device. Yours alone.

— BUILD IT ONCE ON ANY MAC. EVEN INTEL.

Synopsis

Have you ever felt this?

“Wait — didn't I just tell you this?”

The AI on your phone forgot yesterday's conversation. Your family's names, your work, the things keeping you up at night — every time, you start from zero. A small but lonely feeling.

This book is the answer.

I've laid bare every step, every setting — how to build an AI of your own that fixes what cloud AI fundamentally can't: it doesn't remember you, it doesn't know you, and it doesn't keep your life private.

Honestly, about me.

I worked as a network engineer, then as a database engineer, then founded my own company. By 2008, I had stepped away from engineering itself. In 2019, I stepped away from work entirely to raise my children full-time — nearly seven years at home. While I was gone, the world I knew faded: Oracle and Cisco became “legacy,” and AWS and AI took their place.

Even someone like me — eighteen years out of engineering, seven years out of work entirely — could build this. You don't need a new computer. The Mac on your desk right now is enough.

This book is written for two kinds of readers.

If you want to understand how it all works, start from page one. If you'd rather skip the theory and just build the thing, head straight to Udemy and Gumroad. Either path leads to your own AI.

Personal computers gave each of us a machine. Personal AI gives each of us something more — an AI of our own.


One last thing.

Building your own AI feels a little like having a Genie of your own — or, if you grew up in Japan, like having Doraemon. No lamp to rub, no four-dimensional pocket to reach into. But the sense that the world's knowledge is now personal — and yours. An AI that remembers your name, knows your family, and picks up where you left off yesterday.

Intelligence in the cloud. Memory and privacy in your hands.

That's how an AI becomes yours — for life.

Phase 0: Introduction



Throughout this book, the personal AI you build — one that holds your documents, your memory, and your judgment criteria — is called **Your AI**.

The author's own implementation is **AI NOBORu**.

What you build will be yours — **Your AI.**

0-1. Why Build a Personal AI?

Defining “Your AI”

In this book, I call the personal AI you build “**Your AI**.” My own implementation is **AI NOBORu**. Yours will have your own documents, your own memory, your own rules, and your own context. The phrase “Your AI” throughout this book always means the personal AI *you* are building — not AI in general.

A note on “raising” an AI

An AI does not have a personality or consciousness in the human sense. But the documents you show it, the things you have it remember, and the rules you give it shape the direction of its replies, its tone, its judgment, and the kinds of conversations it is good at. We describe that as “the AI growing up.” It is a metaphor, but technically it points at a real fact: the context and configuration you give an AI changes what it tends to output.

What that feeling actually looks like

If you use the Claude or ChatGPT app on your phone every day, that “starting from zero” feeling tends to show up in five specific ways:

- You have to introduce yourself every single time. “I just got back from a work posting.” “I have three children.” “I run a company.” Repeating this gets old fast.
- You want to continue a conversation from a few days ago, but the AI has completely forgotten it.
- You want to hand the AI a PDF, a photo, or your family's schedule directly — but you end up copy-pasting or re-uploading every time.
- You ask, “Which Saturdays are open next month for the family?” — and the AI cannot answer, because it does not know your schedule.
- The fact that your chat history is sitting on someone else's server quietly bothers you.

All of these problems disappear once you build a personal AI of your own.

General-purpose AI vs. an AI of your own

A general-purpose AI is one system fielding questions from hundreds of millions of people. By design, it isn't built to work specifically for you. That's why every conversation starts from scratch, and why it doesn't hold on to your documents.

The system we're building runs on your Mac. It remembers what you've told it, reads your documents, and works with your schedule.




Figure 0-1: A general-purpose AI app (every conversation starts from zero) vs. your own AI (the relationship accumulates over time).

Three values that come from “your own”

Value ①: The relationship accumulates (permanence)

With a general-purpose AI, every conversation starts from zero. With Your AI, the relationship builds. Last week's conversation, last month's decision, the document that arrived three months ago — they stay connected and reachable.

The technical mechanism behind that accumulation is the memory MCP, set up in Phase 4. Of everything in this book, this is the single capability that general-purpose AI apps simply don't offer — and it changes everything about how the AI relates to you over time.

Value ②: You do not have to send your personal data outside

In this book's architecture, your documents, your memory store, and your local working files stay on your Mac. Your Google Calendar and Gmail stay inside your Google account (we never copy them out). When Your AI needs to call a cloud model, only the relevant excerpt needed for that specific task is sent out — not your entire knowledge base.

0-3. The Four Pillars of Permanence

The book's central promise — an AI you can stay with for the long run — comes from four settings working in concert.




Figure 0-2: The four pillars of permanence — System Prompt (constitution), RAG (breadth), Memory Notes (time) ★ most important, and Pinning (focus).

Pillar	Role	In one sentence
System Prompt	The AI's "constitution"	Declare the tone, the judgment criteria, and the rules that must never be broken.
RAG (the bookshelf)	In charge of "breadth"	Make your long-form documents searchable.
Memory Notes (the journal)	In charge of "time" ★	Remember next week, next month, across Workspaces.
Pinning (Pin)	In charge of "focus"	Keep the few documents you can never afford to forget on the desk at all times.

Phase 2 raises three pillars in one stretch: System Prompt, RAG, and Pinning. Phase 4 adds the fourth — Memory. By the end of Phase 4, all four are standing, and this book's central concept is technically

complete.

0-4. What This System Will Be Able to Do

Capability	Phase	Example
Knows what time it is	Phase 3 (time MCP)	"Three years ago today — what was I doing?"
Remembers things	Phase 4 (memory MCP) ★	"That conversation from last week — how did it end up?"
Handles files	Phase 5 (filesystem MCP)	"Read everything in inbox/ and summarize the key points."
Fetches and searches the Web	Phase 6 (fetch + brave)	"What are the latest developments in AI?"
Aggregates a database	Phase 7 (sqlite MCP)	"List the European cities I visited in August 2018, in chronological order."
Reads my calendar	Phase 8 (Google Calendar)	"Tidy up this week's schedule."
Reads my email	Phase 8 (Gmail)	"Which unread emails from this morning are important?"
Looks things up on the shelf	Phase 2 (RAG)	"From last month's meeting notes, pull just the decisions."

Practical examples

Example 1: Morning email triage

"Check today's schedule, then sort the emails I need to reply to this morning by priority. For each one, pull related history from the bookshelf and draft a reply."

- Pulls today's schedule from Google Calendar (Phase 8).
- Looks at the Gmail inbox, surfaces the unanswered messages (Phase 8).
- Searches the bookshelf (RAG) for related history and prior threads (Phase 2).
- Drafts replies in your own writing style, using cues recorded in Memory (Phase 4).
- Saves the result as a Markdown file (Phase 5 filesystem).

4-6. Persistent memory vs. temporary memory

One quality of the memory MCP is worth being clear about: everything you write into it stays forever. That's the point. But it means temporary notes, outdated facts, and mistakes accumulate right alongside everything important.

So this book uses a naming convention: the name you give an entry tells you whether it's permanent or temporary. The author landed on this practice after more than a year of running AI NOBORu.

Category	Naming rule	Expected lifetime	When you delete
Persistent memory	Starts with PERSISTENT_	Years to forever	Basically never
Temporary memory	Starts with TEMP_	Days to weeks	Weekly TEMP_ cleanup

Examples of persistent memory (PERSISTENT_)

```
PERSISTENT_FAMILY_NAMES = RYutaro (11), Yamato (9), Masamune (8)
PERSISTENT_PROFESSIONAL_BACKGROUND = NTT telecom construction → NTT DOCOMO
database → founded a company → full-time homemaker (2019-2026)
PERSISTENT_HEALTH_BASELINE = Annual checkup every April; April 2026 numbers were
fine
PERSISTENT_EDITORIAL_RULES = Educational principle: empathy (“jo”) above all
PERSISTENT_NOBORU_VOICE = Plain declarative English; conclusion first, reasons
after; no filler
```

Examples of temporary memory (TEMP_)

```
TEMP_2026_05_21_CURRENT_PROJECT = Writing Phase 4 of AI of One's Own; Phase 5
starts next week
TEMP_2026_05_21_PENDING_QUESTIONS = Need to confirm bge-m3 chunk limit
TEMP_2026_05_25_AMAZON_ORDER = Ordered laptop stand, delivery 5/28
```

Why use naming conventions at all?

The memory MCP itself has no concept of "temporary" — everything you store is permanent. The PERSISTENT_/TEMP_ convention is how you impose that distinction yourself. Stick to the rule, and your weekly cleanup becomes a single instruction: "delete everything starting with TEMP_." Old material gone; nothing important touched.

Memory naming rule: PERSISTENT_ vs TEMP_

Name everything. PERSISTENT_ lasts for years. TEMP_ gets cleared every week. Memory MCP itself does not distinguish — you do, through naming.




Figure 4-B: PERSISTENT_ entities (teal, left) last for years and are never deleted. TEMP_ entries (amber, right) are cleared in the weekly Sunday cleanup — without touching PERSISTENT_.

4-7. Writing memory in practice

The memory MCP is designed to be driven by Your AI, not by you directly. You could edit the JSON file by hand, but in practice it's far more natural — and far less error-prone — to let Your AI write its own memory.

The basic pattern: write it in via the conversation

```
@agent Please write this to persistent memory:
entity name: PERSISTENT_FAMILY_NAMES
content: RYutaro (11), Yamato (9), Masamune (8)
relation: Noboru Ikuta --[is father of]--> Ryutaro

→ The AI calls the memory MCP's create_entities / create_relations
```

Putting the write rules into the System Prompt

Add the following to the Workspace's System Prompt and Your AI will start volunteering to write things into memory.

```
(An example to add to the System Prompt)
- When an important fact, decision, or rule appears in the conversation,
```

Afterword — Building AI NOBORu, and Coming Home to Myself

On August 9, 2019, I left Japan with my three children — eighteen months, three years and three months, and just turned five. On March 27, 2026, we came home.

For the six years and eight months in between, I was a single father, and raising my children took everything I had.

I set down the work I had done, the career I had built, the place I had held in the world — I let all of it go — and put my children's lives, their education, their growth at the center of everything. It was not a choice I regret. But neither was it a time I would call easy.

Without my noticing, the confidence I had once carried as an engineer, the instincts of a working professional, the conviction of a founder — all of it grew distant, piece by piece. Day after day spent entirely on my children, I had stopped knowing whether I could return to that world. Whether I could build anything again. The questions themselves had begun to dissolve.

As a stay-at-home father, I felt the fear of stepping back into the world — the fear of returning to work — every day, in small ways. And at the same time, something I had only imagined finally became real to me: the fear, and the quiet anxiety beneath it, that people carry when they try to return to working life after time away. For marriage. For childbirth. For raising children. For caring for someone they love. Until then, all I could do was imagine — *that must be hard*. Standing in the same place, I saw the shape of that fear for the first time.

I could have kept going as a stay-at-home father. Nothing in my life required me to force a return. I planned to keep my family at the center of everything, and I quietly assumed I would walk the rest of the way beside my children.

Before I wrote this book, before I built the private AI environment I call AI NOBORu, my confidence had failed me. The more honest way to say it: I had been losing what was left of it.

The design of AI NOBORu began on May 8, 2026. Until that day, I had used ChatGPT and Claude on my iPhone — to look things up, to think out loud, to get help drafting documents, to translate between Japanese and English. That was the whole of it.

I could feel, in my bones, that this intelligence was about to change society. That AI was on its way to becoming the foundation of how the world would work. What I could not yet see was how. How it would actually be used. How it would actually reach into people's lives. Who would carry it into the world. Reading the headlines — trillion-yen investments, tens of trillions — I had quietly accepted that the age of AI would belong to the giants. The rest of us would mostly follow along.

In all honesty, I was standing on the same side of the room as most readers of this book. Someone who used AI, but had not yet found a way to make it their own. Just a stay-at-home father.

Then I finished building AI NOBORu, and everything changed. In the roughly two weeks between May 8, the day I began the design, and May 23, I built not only AI NOBORu but a working AI video studio — an environment that turns out short-form video continuously, mixing my own footage and stills with fully AI-generated material where needed, running on a 2020 Intel-chip iMac.

On top of that, I set up customer support servers for both the Japanese edition, *Sodateru AI*, and the English edition, *AI of One's Own* — for the upcoming Udemy course and the Gumroad product launch — along with an auto-response server for the video studio. All of them run on a 2014 MacBook with sixteen gigabytes of memory.

All of this happened while I remained a stay-at-home father — school runs, lessons, housework, the dozen small things any day asks of you. I was not tucked away in a lab. I built these things inside the rhythm of ordinary life, with AI NOBORu beside me.

How this works day to day — using Claude on my phone, then feeding what matters back into AI NOBORu — is laid out in Phase 9. But none of it would have happened without AI NOBORu. Not one of these things, I am certain, would have taken shape without it.

I had carried a quiet conviction for years. What a person sees, what they hear, what they come to think, the words they leave behind — these should stay with them. Within reach. In a form they can actually use, on the day they need them. Not in a mystical sense. In the most practical sense there is. AI NOBORu is the doorway in.

Once you cross that doorway, things start to get interesting. I did not want to keep the feeling to myself. I wanted to share it with as many people as I could. That impulse became this book. Next, I want to build an AI to help my children with math. Better still: I want to teach them how to build it themselves.

Asking ChatGPT or Claude from your phone is already useful — useful enough that most people stop there. But once you can hand an AI your purposes, your judgment, your sense of what matters, it becomes useful in a much deeper way.

AI does not arrive understanding you. How it behaves depends entirely on what you give it.

What you have it read. What you have it remember. What you tell it to consult every time. What role you give it. From the accumulation of these small choices, the AI slowly becomes an assistant that fits you.

If AI has felt like something difficult — something for other people — I hope this book is, at the very least, a way in.

For the reader who uses ChatGPT or Claude every day, but only ever to ask a question or have a document drafted, and then closes the app.

For the one who feels the age of AI arriving, but cannot yet picture what it will do to daily life, or to work.

For anyone who senses AI will reshape the world, but cannot yet see how.

This book is for you, more than anyone.

And — to the reader preparing to step back into working life after time away, for childbirth, for raising children, for caring for someone you love: I felt that fear in my own body. Your AI can be the first companion at your side as you walk back in.

AI is not only for special people. Not only for companies. Not only for advanced engineers. An individual, holding an AI of their own, to solve the problems that matter to them. An AI that reads your materials, knows your context, and moves toward your purposes. I believe this will become a way of thinking many of us will need.

For me, building AI NOBORu was never only technical work. After the years I had given entirely to raising my children, it was also the work of finding my way back — to the engineer I had been, to the maker I had always been. I think, while I was building the AI, I was quietly rebuilding myself.

My hope is that this book becomes the beginning of an AI environment of your own — and that the AI you build becomes a partner in the problems that matter to you, and grows into a companion you can live alongside for years to come.

AI NOBORu was my first step.

Now it is your turn to build your own.

Noboru Ikuta May 23, 2026

For Yamato, on his tenth birthday.

If You're Already Using AI Tools

Why This Book Still Helps If You Use Claude Code or ChatGPT Projects

This page is an extra note for readers already familiar with Claude Code, ChatGPT Projects, or similar tools.

If you're new to building your own AI environment, feel free to skip ahead to the main content.

Claude Code is a tool for getting work done quickly with AI.

ChatGPT Projects is a room where it's easy to begin working with AI.

Both are excellent AI tools.

But what really matters over the long run is more than knowing how to use any single tool.

Tools change.

Versions change.

Screens change.

The way you use them shifts, bit by bit.

To avoid getting lost each time, you need to understand the structure of the AI environment itself.

This book is about understanding the AI environment itself.

Common pain points for cloud AI users — and how this book helps

Pain point for cloud AI users	What this book helps you understand
Not sure what to put in a Project	RAG = the bookshelf
Not sure what to store in Memory	The kinds of memory and how to use them
Not sure what to write in Claude Code's CLAUDE.md	System Prompt = the AI's constitution
Afraid to let AI touch your files	filesystem = the workshop, with safety rails
Throwing away work logs every time	The idea of feeding logs back into RAG
Get lost when a tool's screen changes	You can navigate it because you understand the structure

How Claude Code, ChatGPT Projects, and AI NOBORu / This Book Differ

Aspect	Claude Code	ChatGPT Projects	AI NOBORu / This Book
Primary purpose	Get work done quickly with AI	Make it easy to start working with AI	Understand the AI environment itself
Reader's	Use a handy tool	Use a ready-made room	Build the AI's structure

Aspect	Claude Code	ChatGPT Projects	AI NOBORu / This Book
experience			yourself
What you learn	Instructions, edits, execution, permissions	Instructions, file organization, continuing conversations	System Prompt / RAG / Memory / MCP / filesystem / SQLite / backup
Outcome	Work moves faster	You can work with all your context in one place	Both the AI and the user grow
Strength	Ready to go from day one	An easy entry point	You build long-term AI literacy
Weakness	You can use it without understanding the mechanics	Much of the mechanics live on the service side	Setup takes effort

This book doesn't treat AI as just a chat tool.

System Prompt is "the AI's constitution."

RAG is "the bookshelf."

Memory is the AI's long-term memory.

The filesystem is the AI's workspace.

MCP is the toolbox.

Backup is insurance for the long run.

What do you let the AI read?

What do you want it to remember?

How far do you let it work?

Which tools do you give it?

How do you protect it, and how do you back it up?

As you work through each of these, you yourself come to understand how your AI environment is built.

What These Concepts Look Like in Claude Code and ChatGPT Projects

This book's term	Technical term	In Claude Code	In ChatGPT Projects
The AI's constitution	System Prompt / Instruction	CLAUDE.md / project rules	Project instructions
The bookshelf	RAG / Knowledge Base	Markdown, README, and reference files inside the project	Project files

This book's term	Technical term	In Claude Code	In ChatGPT Projects
Memory	Memory / Persistent Memory	Auto Memory / memory files	Project memory / Saved memory
The workshop	Workspace / filesystem	Project folder / working directory	Project / Canvas / uploaded workspace
Tools	MCP / Tools	Tools / MCP servers	Apps / Connectors
Safety rails	permissions / allowed paths	Allow / Ask / Deny / hooks	Connector permissions / user confirmation
Work log	chat log / execution log	session transcript / command history	Project chats
Backup	backup / export	Git / local backup / session export	Saving originals / export

Just as raising a child also raises the parent, raising Your AI is itself the shortest path to understanding how your AI environment is built.

This book is not a manual for Claude Code.

Nor is it just a guide to using ChatGPT Projects.

This is a book about using AI — and about becoming someone who can raise an AI.

INDEX

Synopsis	I
If You're Already Using AI Tools	II
Phase 0: Introduction	1
0-1. Why Build a Personal AI?	3
0-2. How Each Setting Changes the AI	5
0-3. The Four Pillars of Permanence	7
0-4. What This System Will Be Able to Do	8
0-5. The Big Picture	10
0-6. The Map of the Book (Nine Phases + Appendix)	11
0-7. What You Will Need (Time, Money, Skills)	12
0-8. How to Read This Book	13
0-9. The Author's Build Environment	15
0-10. Setup Samples (a Preview)	15
Summary of Phase 0	16
Phase 1: Installing AnythingLLM (Docker Edition)	17
1-1. The metaphor — secure your land before you build the house	19
1-2. From here on, you'll be using your hands	20
1-3. Phase 1 — wrap-up	21
Phase 2: Giving Your AI a Personality, a Brain, and Knowledge	22
📖 Terminology used in this book	24
2-1. The seven skills you'll equip the workshop with	25
2-2. The metaphor — moving the residents, the bookshelves, and the constitution into the workshop	27
2-3. How this chapter fits the book's concept	27
2-4. Designing and creating Workspaces	27
2-5. Connecting the LLM Provider (the brain)	29
2-6. The Agent Provider (the hand that uses the tools)	31
2-7. The Embedding Engine — Ollama + bge-m3	31
2-8. Writing the System Prompt (the AI's Constitution)	32
2-9. Loading documents into RAG (the Bookshelf)	34
2-10. Configuring Pinning	35
2-11. Verifying — are three of the four pillars actually up?	37
2-12. Common stumbles and how to handle them	37
2-13. How this chapter relates to the concept — a recap	38
2-14. Completion checklist	39
[A look at AI NOBORu] — The author's actual personality and knowledge setup	40
Phase 3: Giving Your AI a Sense of Time (time MCP)	42
3-1. The metaphor — handing Your AI a clock	43
3-2. How this chapter fits the book's concept	44
3-3. A short note on what MCP is	45
3-4. Initializing the MCP configuration file	46
3-5. Adding the time MCP	46
3-6. UI Refresh and verifying the server is up	47
3-7. Testing time fetch from Your AI	48
3-8. Common stumbles and how to handle them	49
3-9. How this chapter relates to the concept — a recap	50
3-10. Completion checklist	50
[A look at AI NOBORu] — how a sense of time is actually being used	51
Phase 4: Giving Your AI a Memory (memory MCP)	52
4-1. The metaphor — handing Your AI a journal (and it's not an ordinary one)	53
4-2. How this chapter fits the concept — why this is the heart	54
4-3. What the memory MCP is — the Knowledge Graph approach	55

4-4. Adding the memory MCP	56
4-5. Verifying the memory MCP is up	57
4-6. Persistent memory vs. temporary memory	58
4-7. Writing memory in practice	59
4-8. Searching and reading memory	60
4-9. Maintaining memory (weekly)	61
4-10. Common stumbles and how to handle them	61
4-11. How this chapter relates to the concept — a recap	62
4-12. Completion checklist	63
[A look at AI NOBORu] — the author's actual Memory practice	64
Phase 5: Letting Your AI Hold, Update, and Save Your Documents	65
5-1. What the filesystem lets you do	67
5-2. How this chapter fits the concept	69
5-3. Adding the filesystem MCP	70
5-4. Starting it up and verifying	72
5-5. Common stumbles and how to handle them	73
5-6. Security practices for daily use	73
5-7. How this chapter relates to the concept — a recap	73
5-8. Completion checklist	74
[A look at AI NOBORu] — the author's filesystem practice	74
Phase 6: Giving Your AI the Internet (fetch + brave-search MCP)	76
6-1. The metaphor — giving Your AI eyes	78
6-2. How this chapter fits the concept	78
6-3. Adding the fetch MCP	79
6-4. Adding the brave-search MCP	80
6-5. When to use which — fetch vs. brave-search	81
6-6. Common stumbles and how to handle them	82
6-7. How this chapter relates to the concept — a recap	83
6-8. Completion checklist	83
[A look at AI NOBORu] — the author's web-integration practice	83
Phase 7: Giving Your AI a Database (sqlite MCP)	85
7-1. The metaphor — handing Your AI a database	86
7-2. How this chapter fits the concept	87
7-3. Adding the sqlite MCP	88
7-4. Table design — what to put in the database	88
7-5. Quick test — let the AI write the SQL	89
7-6. Common stumbles and how to handle them	91
7-7. How this chapter relates to the concept — a recap	91
7-8. Completion checklist	91
[A look at AI NOBORu] — the author's sqlite practice	92
Phase 8: Connecting Google Calendar and Gmail	93
8-1. The metaphor — connecting Your AI to real life	95
8-2. How this chapter fits the concept	95
8-3. Choosing the connection method — why Apps Script	95
8-4. Creating the Apps Script project	96
8-5. Deploying as a Web App	99
8-6. Configuring the call from AnythingLLM	99
8-7. Common stumbles and how to handle them	101
8-8. How this chapter relates to the concept — a recap	102
8-9. Completion checklist	102
[A look at AI NOBORu] — the author's Calendar / Gmail practice	103
Phase 9: Cost Optimization and Operations	104
9-1. The metaphor — let Your AI pick a brain that fits the moment	106
9-2. Where are you now? — what you're actually spending	106
9-3. Designing the model split	106
9-3-1. How to grow Your AI for “almost free”	108

9-4. The six cost-reduction principles	109
9-5. Monthly monitoring routine	111
9-5-1. AI NOBORu in practice — how the author manages cost.....	112
9-5-2. The design philosophy — “decide on a lifetime cost first”	113
9-6. Common stumbles and how to handle them	114
9-7. How this chapter relates to the concept — a recap.....	114
9-8. Completion checklist.....	115
Appendix A: Backup & Recovery	116
A-1. Why This Chapter Exists	116
A-2. What to Protect — Priority Table.....	116
A-3. Three-Tier Backup Design	117
A-4. Weekly Backup Script	117
A-5. Monthly Backup Script	118
A-6. TEMP_ Cleanup Routine	119
A-7. Restore Procedures	120
A-8. Troubleshooting	121
A-9. Completion Checklist	121
A-10. Data Protection Policy.....	122
A-11. AI NOBORu Backup Setup	123
Appendix B: Docker Setup (Phase 1 — Detailed Instructions).....	125
B-1. Identify Your Mac	125
B-2. Install Docker Desktop	125
B-3. Create the Working Directories	126
B-3b. .env File (Optional).....	126
B-4. Build the Custom Docker Image	127
B-5. Create docker-compose.yml	127
B-6. Start the Container	128
B-7. Verification Checklist.....	129
B-8. Troubleshooting	129
B-9. How This Fits Into the Book	129
B-10. AI NOBORu Reference Configuration.....	130
Appendix C: Complete Configuration Files	131
C-1. anythingllm_mcp_servers.json — Full MCP Configuration.....	131
C-2. Google Apps Script — Complete Code.gs.....	132
C-3. SQLite Initialization SQL	135
Appendix D: Reaching Your Home AI Server Safely with Tailscale	137
D-1. Why This Appendix Exists	137
D-2. Why Tailscale Instead of Opening a Port.....	137
D-3. How Tailscale Works (Briefly)	138
D-4. The Overall Picture	138
D-5. Preflight Check (Don't Skip This)	139
D-6. Install Tailscale on the iMac.....	140
D-7. Install Tailscale on the iPhone	140
D-8. Find the iMac's Tailscale IP	141
D-9. Connect to AnythingLLM From Your iPhone	141
D-10. Connect From a MacBook	141
D-11. MagicDNS (Optional)	141
D-12. Security Notes.....	142
D-13. Troubleshooting	142
D-14. Notes on Real-World Operation.....	143
D-15. In Closing.....	144
Afterword — Building AI NOBORu, and Coming Home to Myself	145
About the Author	148